

# Interview Review Chart

Karl Solomon

December 25, 2024

## 1 Parallelism

### 1.1 Concepts

**Concurrency** When two or more tasks can start/run/complete in overlapping time periods. This does not necessarily mean they'll be running in overlapping time periods. Examples include:  
RTOS

**Parallelism** When tasks run at the same time (e.g. on a multi-core CPU)

**Multithreading** When multiple tasks are running on a CPU. This can be implemented truly parallel where each task has access to separate HW/core. However, more common in desktop applications is SMT.

**SMT (Simultaneous Multithreading)** Multiple threads share 1 core. The thread instructions are pipelines s.t. they run mostly in-parallel, and when one is waiting for I/O the other can run uninhibited, however since only one thread can access a dedicated HW block at any given time they are not truly parallel.

### 1.2 Implementation

- Concurrency and Multithreading

- Basic concepts of threads and processes
- Thread synchronization mechanisms (mutexes, semaphores, locks)
- Race conditions:
  1. Multiple threads accessing a shared resource
  2. At least one thread writes to the resource
  3. Lack of proper synchronization
- deadlocks

```
1 std::mutex m1, m2;  
2 // Must have 2 locks where they are locked in different orders  
  in different locations/threads  
3 thread1: m1.lock(); m2.lock(); m1.unlock(); m2.unlock();  
4 thread2: m2.lock(); m1.lock(); m2.unlock(); m1.unlock();
```

- Atomic operations

```
1 std::atomic<T> var; // where T is a primitive type  
2 var = val;  
3 var.load(val);  
4 var.store(val);  
5 var.wait(); // waits until the value changes
```

```

6 T current = var.exchange(val); // writes val to var and gets
  previous/current value of var
7 // Compare-and-swap. Atomically compares object.value with
  that of expected. If bitwise-equal then replaces former
  with desired. Otherwise loads actual value into expected (
  via load operation)
8 bool res = var.compare_exchange_strong(expected, desired); //
  preferred when don't expect high contention and cost of
  retrying is significant. Simpler, but slower.
9 bool res = var.compare_exchange_weak(expected, desired); //
  preferred if you're anyways retrying in a loop or cost of
  retrying is low. More efficient, but more complex.

```

```

1 // Weak usage
2 std::atomic<int> value{0};
3 int expected = 0;
4 int desired = 1;
5 while(!value.compare_exchange_weak(expected, desired)){
6     ; // handle spurious failures
7 }
8
9 // Strong usage
10 if(value.compare_exchange_strong(expected, desired)){
11     // op successful
12 }
13 else
14 {
15     // op failed
16 }

```

- Thread-Safe Data Structures:

- Concurrent collections
- Concurrent collections (e.g., ConcurrentQueue, ConcurrentBag)
- Lock-free data structures
- Understanding the differences between thread-safe and non-thread-safe collections

- Design Patterns for Concurrency:

- Producer-Consumer pattern
- Readers-Writer pattern
- Thread pool pattern

- Language-Specific Concurrency Features for C++:

- std::thread
- <atomic>

- Callback Mechanisms:

- Function pointers

```

1 #include <stdio.h>
2
3 // Define a struct with function pointers for arithmetic
  operations
4 typedef struct {
5     int (*add)(int, int);
6 } ArithmeticOperations;
7

```

```

8 // Define the functions for arithmetic operations
9 int add(int a, int b) { return a + b; }
10
11 int main() {
12     // Initialize the struct with function pointers
13     ArithmeticOperations ops;
14     ops.add = add;
15     // Use the function pointers to perform operations
16     int x = 10, y = 5;
17     printf("Add: %d + %d = %d\n", x, y, ops.add(x, y));
18     return 0;
19 }

```

- Delegates (in languages that support them)
- Lambda expressions

- Performance Considerations:

- Understanding the overhead of different synchronization mechanisms  
Generally best practice to measure performance in single-threaded vs multithreaded/parallel environments. Since there is overhead with creating/cleaning up threads/processes it can make your program run slower in smaller data sets.
- Balancing thread safety with performance

- Testing Multithreaded Code:

- Techniques for writing unit tests for concurrent code
- Tools for detecting race conditions and deadlocks
  - \* **Helgrind:** Part of Valgrind suite. Checks for race conditions, but slow.
  - \* **ThreadSanitizer:** Compiler flag in llvm/clang. Faster than Helgrind.
  - \* **RacerD:** Meta's C++-specific concurrent static analyzer. Good for target code-bases.
  - \* **Clang Static Analyzer:** Detects some simple conditions.

- Distributed Systems Concepts:

While not directly related to this problem, understanding concepts like eventual consistency and distributed locking can be beneficial

- Algorithms for Concurrent Operations:

- Compare-And-Swap (CAS) operations
- Lock-free algorithms

- Memory Models:

- Understanding memory barriers and volatile variables
- Cache coherence issues in multi-core systems

**Cache Coherence** The process of ensuring that data is stored in multiple caches within a multiprocessor system is consistent and synchronized.

This ensures that all processors have a *consistent* view of shared memory. Cache coherence protocols manage the flow of data between caches, updating cache lines and tracking the status of shared data. This can be complicated because it requires balancing performance and coherence overhead. The 2 main types of protocols are Directory-based and Snoop-based.

**Directory-based** The sharing status of a block of physical memory is kept in just one location (the directory). The directory can also be distributed to improve scalability. Communication is established using point-to-point requests through the interconnection network.

**Snoop-based** Every cache that has a copy of the data from a block of physical memory also has a copy of the sharing status of the block, but no centralized state is kept. Caches are all accessible via some broadcast medium (a bus or switch), and all cache controllers monitor or *snoop* on the medium to determine whether or not they have a copy of a block that is requested on a bus or switch access. Requires broadcast, since caching information is at processors. This is useful for small-scale machines.

**Point of Coherency (PoC)** Point at which all agents in a system which can access memory are guaranteed to see the same data.

**Migration** Data is migrated to the local cache levels.

**Replication** The same data is replicated across all caches.

Assume Snoop-based protocol. There are 2 ways to maintain coherence:

1. **Write Invalidate Protocol:** Ensure that a processor has exclusive access to a data item before it writes that item. This is most common protocol.
2. **Write Broadcast/Update:** All cached copies are updated simultaneously. This requires more bandwidth. When multiple updates happen to the same location, unnecessary updates are done. However, this is a lower latency between write/read.

- Practice Problems:

- Implement a thread-safe singleton

```
1 #include <iostream>
2 #include <mutex>
3
4 class Singleton {
5 public:
6     // Delete copy constructor and assignment operator to
        // prevent creation of additional instances through
        // copying
7     Singleton(const Singleton&) = delete;
8     Singleton& operator=(const Singleton&) = delete;
9
10    // Static method to get the instance of the singleton
11    static Singleton& getInstance() {
12        static Singleton instance; // Static Local variables
        // are guaranteed to be thread-safe in C++11 and later
13        return instance;
14    }
15
16    void showMessage() { std::cout << "Singleton instance
        accessed!" << std::endl; }
17
18 private:
19    // Private constructor to prevent instantiation
20    Singleton() { std::cout << "Singleton instance created!"
        << std::endl; }
21 };
22
23 int main() {
24     // Access the singleton instance
25     Singleton& singleton = Singleton::getInstance();
26     singleton.showMessage();
```

```

27
28     return 0;
29 }

```

– Create a simple producer-consumer queue

```

1  #include <chrono>
2  #include <condition_variable>
3  #include <iostream>
4  #include <mutex>
5  #include <queue>
6  #include <thread>
7
8  class ThreadSafeQueue {
9  public:
10     void enqueue(int item) {
11         std::lock_guard<std::mutex> lock(mutex_);
12         queue_.push(item);
13         cond_var_.notify_one();
14     }
15
16     int dequeue() {
17         std::unique_lock<std::mutex> lock(mutex_);
18         cond_var_.wait(lock, [this] { return !queue_.empty(); });
19         int item = queue_.front();
20         queue_.pop();
21         return item;
22     }
23
24 private:
25     std::queue<int> queue_;
26     std::mutex mutex_;
27     std::condition_variable cond_var_;
28 };
29
30 void producer(ThreadSafeQueue& queue, int numItems) {
31     for (int i = 0; i < numItems; ++i) {
32         std::this_thread::sleep_for(std::chrono::milliseconds
33             (100)); // Simulate work
34         queue.enqueue(i);
35         std::cout << "Produced: " << i << std::endl;
36     }
37 }
38
39 void consumer(ThreadSafeQueue& queue, int numItems) {
40     for (int i = 0; i < numItems; ++i) {
41         int item = queue.dequeue();
42         std::cout << "Consumed: " << item << std::endl;
43     }
44 }
45
46 int main() {
47     ThreadSafeQueue queue;
48     const int numItems = 10;
49
50     std::thread producerThread(producer, std::ref(queue),
51         numItems);
52     std::thread consumerThread(consumer, std::ref(queue),
53         numItems);
54
55     producerThread.join();
56     consumerThread.join();

```

```

55     return 0;
56 }

```

– Implement a basic thread pool

```

1  #include <execution>    // Required for std::execution::seq
2  #include <iostream>
3  #include <numeric>
4  #include <random>
5  #include <thread>
6  #include <vector>
7
8  void accumulateRandomNumbers(int threadID) {
9      static std::mutex m;
10     std::random_device rd;
11     std::mt19937 gen(rd());
12     std::uniform_real_distribution<double> dis(0.0, 1.0);
13
14     std::vector<double> numbers(1024 * 1024);
15     for (auto& num : numbers) {
16         num = dis(gen);
17     }
18
19     double sum = std::reduce(std::execution::seq, numbers.
20         begin(), numbers.end());
21     std::lock_guard<std::mutex> lock(m);
22     std::cout << "Thread " << threadID << " accumulated sum: "
23         << sum << std::endl;
24 }
25
26 int main() {
27     const int numThreads = std::thread::hardware_concurrency();
28     ; // Number of threads in the pool
29     std::vector<std::thread> threadPool;
30
31     // Create and launch threads
32     for (int i = 0; i < numThreads; ++i) {
33         threadPool.emplace_back(accumulateRandomNumbers, i);
34     }
35
36     // Join threads to the main thread
37     for (auto& thread : threadPool) {
38         thread.join();
39     }
40
41     return 0;
42 }

```

– Solve classic concurrency problems like the dining philosophers problem

```

1  #include <chrono>
2  #include <condition_variable>
3  #include <iostream>
4  #include <mutex>
5  #include <thread>
6  #include <vector>
7
8  class DiningPhilosophers {
9  public:
10     DiningPhilosophers(int numPhilosophers)
11         : numPhilosophers_(numPhilosophers), states_(
12             numPhilosophers, State::THINKING) {}

```

```

13     void philosopher(int id) {
14         while (true) {
15             think(id);
16             pickUpForks(id);
17             eat(id);
18             putDownForks(id);
19         }
20     }
21
22 private:
23     enum class State { THINKING, HUNGRY, EATING };
24
25     void think(int id) {
26         std::cout << "Philosopher " << id << " is thinking."
27             << std::endl;
28         std::this_thread::sleep_for(std::chrono::milliseconds
29             (1000));
30     }
31
32     void eat(int id) {
33         std::cout << "Philosopher " << id << " is eating." <<
34             std::endl;
35         std::this_thread::sleep_for(std::chrono::milliseconds
36             (1000));
37     }
38
39     void pickUpForks(int id) {
40         std::unique_lock<std::mutex> lock(mutex_);
41         states_[id] = State::HUNGRY;
42         cond_var_.wait(lock, [this, id] { return canEat(id);
43             });
44         states_[id] = State::EATING;
45     }
46
47     void putDownForks(int id) {
48         std::unique_lock<std::mutex> lock(mutex_);
49         states_[id] = State::THINKING;
50         cond_var_.notify_all();
51     }
52
53     bool canEat(int id) {
54         int left = (id + numPhilosophers_ - 1) %
55             numPhilosophers_;
56         int right = (id + 1) % numPhilosophers_;
57         return states_[id] == State::HUNGRY && states_[left]
58             != State::EATING && states_[right] != State::EATING
59             ;
60     }
61
62     int numPhilosophers_;
63     std::vector<State> states_;
64     std::mutex mutex_;
65     std::condition_variable cond_var_;
66 };
67
68 int main() {
69     const int numPhilosophers = 5;
70     DiningPhilosophers diningPhilosophers(numPhilosophers);
71     std::vector<std::thread> threads;
72
73     for (int i = 0; i < numPhilosophers; ++i) {
74         threads.emplace_back(&DiningPhilosophers::philosopher,
75             &diningPhilosophers, i);
76     }
77 }

```

```

67     }
68
69     for (auto& thread : threads) {
70         thread.join();
71     }
72
73     return 0;
74 }

```

Write a few test cases in addition to the solution. Remember, for interviews, it's not just about knowing the solutions, but also being able to explain your reasoning, discuss trade-offs, and analyze the performance and correctness of your solutions. Lastly, be prepared to write code on a whiteboard or in a simple text editor. Practice implementing these concepts without relying on an IDE's features.

## 2 Operating Systems

### 2.1 Caches

**TLB** The TLB is a small, fast, and fast-access memory that is used to translate virtual memory addresses into physical memory addresses. a.k.a. Translation Lookaside Buffer.

### 2.2 Basic Concepts

- **TLB: Translate Lookaside Buffer**
- **Processes and Threads**
  - Process creation and termination
  - Thread lifecycle and management
- **Memory Management**
  - Virtual memory
  - Paging and segmentation
- **File Systems**
  - File system structure
  - File operations and permissions

## 3 C

### 3.1 Preprocessor

```

1 # // stringizes the macro parameter
2 #define stringify(x) #x
3 #define foo 1
4 stringify(foo) // --> evaluates to "foo", NOT "1"

```

```

1 ## // concatenates the macro parameter
2 #define COMMAND(NAME) {#NAME, NAME ## _command}
3 struct command commands[] = {
4     COMMAND(quit), // equivalent to {quit_command}
5     COMMAND(help), // equivalent to {help_command}
6 }

```



- predefined macros
  - `__FILE__`
  - `__LINE__`
  - `__DATE__`
  - `__TIME__`
  - `__STDC_VERSION__`
  - `__cplusplus`
- item2
- item3
- item4

## 3.2 Peripherals

### • I2C

SDA is data, SCL is clock. PURs typically in the 1-4.7k range. Too weak = slow comm and errors. Clocks are usually 100k-1MHz. Addr can be 7 or 10 bit. This is rate-limiter for number of slaves, though line impedance would increase for each slave. Here are some use usage examples:

1. Master sends START and slave Addr
2. Master sends data to slave
3. Master terminates with a STOP

1. Master sends START and slave Addr
2. Master sends data to slave
3. Master sends repeatedSTART and either sends more data to slave or receives data from slave.
4. Master sends STOP

### • SPI

Serial Peripheral Interface (SPI) is a synchronous serial communication protocol used for short-distance communication, primarily in embedded systems. It uses a master-slave architecture with a single master and multiple slaves. Communication is full-duplex, and it requires four wires: MOSI, MISO, SCLK, and SS.

### • UART

Universal Asynchronous Receiver-Transmitter (UART) is a hardware communication protocol that uses asynchronous serial communication with configurable speed. It is commonly used for communication between microcontrollers and peripherals. UART requires only two wires: TX (transmit) and RX (receive).

### • USB

Universal Serial Bus (USB) is an industry-standard for short-distance digital data communications. It supports plug-and-play installation and hot swapping. USB is used for connecting peripherals such as keyboards, mice, printers, and external storage devices to computers.

### • HDMI

High-Definition Multimedia Interface (HDMI) is a proprietary audio/video interface for transmitting uncompressed video data and compressed or uncompressed digital audio data from an HDMI-compliant source device to a compatible display device. It is commonly used for connecting devices like TVs, monitors, and projectors.

## 4 C++

### 4.1 <algorithm>

- **batchOperations**

– `for_each`, `ranges::for_each`, `for_each_n`, `ranges::for_each_n`

```
1 std::vector<int> numbers = {1, 2, 3, 4, 5};
2
3 // Use std::for_each to print each element
4 std::for_each(numbers.begin(), numbers.end(), [](int n) { std
  ::cout << n << " "; });
5
6 // Use std::for_each_n to print the first 3 elements
7 std::for_each_n(numbers.begin(), 3, [](int n) { std::cout << n
  << " "; });
8
9 // Use std::ranges::for_each to print each element
10 std::ranges::for_each(numbers, [](int n) { std::cout << n << "
  "; });
11
12 // Use std::ranges::for_each_n to print the first 3 elements
13 std::ranges::for_each_n(numbers.begin(), 3, [](int n) { std::
  cout << n << " "; });
```

- **Search Operations**

– `all_of`, `any_of`, `none_of`

```
1 std::vector<int> numbers = {1, 2, 3, 4, 5};
2 // Use std::all_of to check if all elements are positive
3 bool allPositive = std::all_of(numbers.begin(), numbers.end(),
  [](int n) { return n > 0; });
4 // Use std::any_of to check if any element is greater than 4
5 bool anyGreaterThanFour = std::any_of(numbers.begin(), numbers
  .end(), [](int n) { return n > 4; });
6 // Use std::none_of to check if no elements are negative
7 bool noneNegative = std::none_of(numbers.begin(), numbers.end
  (), [](int n) { return n < 0; });
```

– `ranges::contains`, `ranges::contains_subrange`

```
1 std::vector<int> numbers = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};
2
3 // Check if the range contains the value 5
4 bool containsFive = std::ranges::contains(numbers, 5);
5 // Check if the range contains the value 11
6 bool containsEleven = std::ranges::contains(numbers, 11);
7 // Define a subrange to check
8 std::vector<int> subrange = {4, 5, 6};
9 // Check if the range contains the subrange
10 bool containsSubrange = std::ranges::contains_subrange(numbers
  , subrange);
11 // Define another subrange to check
12 std::vector<int> nonExistentSubrange = {7, 8, 11};
13 // Check if the range contains the non-existent subrange
14 bool containsNonExistentSubrange = std::ranges::
  contains_subrange(numbers, nonExistentSubrange);
```

- find, find\_if, find\_if\_not, ranges::find, ranges::find\_if, ranges::find\_if\_not

```
1 #include <algorithm>
2 #include <iostream>
3 #include <ranges>
4 #include <vector>
5
6 int main() {
7     std::vector<int> numbers = {1, 2, 3, 4, 5, 6, 7, 8, 9,
8                               10};
9
10    // Use std::find to find the first occurrence of 5
11    auto it = std::find(numbers.begin(), numbers.end(), 5);
12
13    // find the first even number
14    it = std::find_if(numbers.begin(), numbers.end(), [](int n
15    ) { return n % 2 == 0; });
16    // find the first odd number
17    it = std::find_if_not(numbers.begin(), numbers.end(), [](
18    int n) { return n % 2 == 0; });
19    // find the first occurrence of 5
20    auto range_it = std::ranges::find(numbers, 5);
21    // find the first even number
22    range_it = std::ranges::find_if(numbers, [](int n) {
23    return n % 2 == 0; });
24    // find the first odd number
25    range_it = std::ranges::find_if_not(numbers, [](int n) {
26    return n % 2 == 0; });
27 }
```

- find\_last, find\_last\_if, find\_last\_if\_not, find\_end, ranges::find\_end

```
1 #include <algorithm>
2 #include <iostream>
3 #include <ranges>
4 #include <vector>
5
6 std::vector<int> numbers = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 5,
7                          6, 7};
8
9 // find the last occurrence of 5
10 auto lastFive = std::ranges::find_last(numbers, 5);
11 // find the last even number
12 auto lastEven = std::ranges::find_last_if(numbers, [](int n) {
13 return n % 2 == 0; });
14 // find the last odd number
15 auto lastOdd = std::ranges::find_last_if_not(numbers, [](int n
16 ) { return n % 2 == 0; });
17
18 // Define a subrange to find
19 std::vector<int> subrange = {5, 6, 7};
20
21 // find the last occurrence of the subrange
22 auto lastSubrange = std::find_end(numbers.begin(), numbers.end
23 (), subrange.begin(), subrange.end());
24
25 // find the last occurrence of the subrange
26 auto lastSubrangeRange = std::ranges::find_end(numbers,
27 subrange);
```

• find\_end, ranges::find\_end, find\_first\_of, ranges::find\_first\_of



## 4.2 Classes

- **Class Definition**

- Syntax and structure
- Access specifiers: `public`, `private`, `protected`

- **Inheritance**

- Single and multiple inheritance
- Virtual inheritance

- **Polymorphism**

- Function overloading
- Operator overloading
- Virtual functions and abstract classes

## 4.3 Containers

- Sequence

- **array**

```
1 std::array<int, 3> arr; // uninitialized (whatever was in
   memory before)
2 std::array<int, 3> arr = {}; // initialized as 0s
3 std::array<int, 3> arr1 = {1, 2, 3};
4 std::array<int, 3> arr2{1, 2, 4};
5 arr1.fill(0); // fills array with 0s
6 arr1.swap(arr2); // swaps contents of arr1 and arr2
```

- **vector**

```
1 std::vector<int> v;
2 v.capacity(); // size of currently allocated memory
3 v.shrink_to_fit(); // releases unused memory
4 v.reserve(100); // pre-allocates 100 elements
5 v.clear(); // erases all elements
6 v.erase(v.begin()); // erases first element
7 v.push_back(1); // adds 1 to the end
8 v.rbegin(); // reverses iterator
9 std::erase_if(v, [](int x) { return x > 10; }); // removes all
   elements > 10
10 std::vector<Pair<int, int>> classV;
11 classV.emplace_back(10, 1); // create Pair object and push to
   back
```

- **inplace\_vector**

- **deque**

- Associative

- **Set**
- **Map**
- **Multiset**
- **Multimap**

- Unordered Associative

- **unordered\_set**
- **unordered\_map**

- `unordered_multiset`
- `unordered_multimap`

- **Adaptors**

- `stack`
- `queue`
- `priority_queue`
- `flat_set`
- `flat_map`
- `flat_multiset`
- `flat_multimap`

## 4.4 Modern C++

- C++11

- **Alias Templates**

```

1 #include <iostream>
2 #include <map>
3 #include <string>
4 #include <vector>
5
6 // Alias template for a vector of a specific type
7 template <typename T>
8 using Vector = std::vector<T>;
9
10 // Alias template for a map with string keys and a specific
    value type
11 template <typename V>
12 using StringMap = std::map<std::string, V>;
13
14 int main() {
15     // Using the alias template for a vector of integers
16     Vector<int> intVector = {1, 2, 3, 4, 5};
17     std::cout << "Vector of integers: ";
18     for (const auto& elem : intVector) {
19         std::cout << elem << " ";
20     }
21     std::cout << std::endl;
22
23     // Using the alias template for a map with string keys and
        integer values
24     StringMap<int> ageMap = {{ "Alice", 30}, {"Bob", 25}, {"
        Charlie", 35}};
25     std::cout << "Map of ages: ";
26     for (const auto& pair : ageMap) {
27         std::cout << pair.first << ": " << pair.second << " ";
28     }
29     std::cout << std::endl;
30
31     return 0;
32 }

```

- **atomic**

Well-defined behavior in the event of RMW race condition. Accesses to atomics may establish inter-thread synchronization and order non-atomic accesses.

```

1 atomic_bool b; // same as std::atomic<bool> b;

```

- **auto**
- **constexpr**
- **final**

- \* Specifies that a class cannot be inherited from.
- \* When used in a virtual function, specifies that the function cannot be overridden by a derived class.
- \* **final** is also a legal variable/function name. Only has special meaning in member function declaration or class head.

```

1 struct Base
2 {
3     virtual void foo();
4 };
5 struct A : Base
6 {
7     void foo() final; // Base::foo is overridden and A::foo is
                        // the final override
8     void bar() final; // Error: bar cannot be final as it is
                        // non-virtual
9 };
10
11 struct B final : A // struct B is final
12 {
13     void foo() override; // Error: foo cannot be overridden as
                            // it is final in A
14 };
15
16 struct C : B {}; // Error: B is final

```

- **initializer list**

```

1  /*
2   * In this program:
3   *   Vector Initialization: A 'std::vector' is initialized
4   *   using an initializer list, which provides a concise way to
5   *   initialize containers with a list of values. Class
6   *   Constructor: The 'MyClass' constructor takes an
7   *   'std::initializer_list<int>' as a parameter, allowing
8   *   objects of 'MyClass' to be initialized with a list of
9   *   integers.
10  *   Function Parameter: The 'printList' function takes an 'std
11  *   ::initializer_list<std::string>' as a parameter,
12  *   demonstrating how initializer lists can be used to pass a
13  *   variable number of arguments to a function. This program
14  *   demonstrates the flexibility and convenience of using
15  *   initializer lists in various contexts in C++.
16  */
17
18 #include <initializer_list>
19 #include <iostream>
20 #include <vector>
21
22 class MyClass {
23 public:
24     MyClass(std::initializer_list<int> list) {
25         for (auto elem : list) {
26             data_.push_back(elem);
27         }
28     }
29
30     void print() const {

```

```

24         for (auto elem : data_) {
25             std::cout << elem << " ";
26         }
27         std::cout << std::endl;
28     }
29
30 private:
31     std::vector<int> data_;
32 };
33
34 void printList(std::initializer_list<std::string> list) {
35     for (const auto& elem : list) {
36         std::cout << elem << " ";
37     }
38     std::cout << std::endl;
39 }
40
41 int main() {
42     // Initializing a vector using an initializer list
43     std::vector<int> vec = {1, 2, 3, 4, 5};
44     std::cout << "Vector elements: ";
45     for (int v : vec) {
46         std::cout << v << " ";
47     }
48     std::cout << std::endl;
49
50     // Using initializer list in a class constructor
51     MyClass myObject = {10, 20, 30, 40, 50};
52     std::cout << "MyClass elements: ";
53     myObject.print();
54
55     // Passing an initializer list to a function
56     std::cout << "String list: ";
57     printList({"Hello", "World", "from", "initializer", "list"});
58
59     return 0;
60 }

```

#### – iota

```

1 void iota(ForwardIterator begin, ForwardIterator end, T v); //
    fills range [first-last] with sequentially increasing
    values starting at v in begin

```

#### – lambdas

```

1 #include <algorithm>
2 #include <functional>
3 #include <iostream>
4 #include <vector>
5
6 int main() {
7     // Basic lambda with no capture
8     auto greet = []() { std::cout << "Hello, World!" << std::
9         endl; };
10    greet();
11
12    // Lambda with capture by value
13    int a = 10;
14    auto captureByValue = [a]() { std::cout << "Captured by
15        value: " << a << std::endl; };
16    captureByValue();

```



```

16 // Lambda with capture by reference
17 int b = 20;
18 auto captureByReference = [&b]() {
19     b += 10;
20     std::cout << "Captured by reference: " << b << std::
        endl;
21 };
22 captureByReference();
23 std::cout << "Modified b: " << b << std::endl;
24
25 // Lambda with explicit return type
26 auto add = [](int x, int y) -> int { return x + y; };
27 std::cout << "Sum: " << add(3, 4) << std::endl;
28
29 // Generic lambda
30 auto multiply = [](auto x, auto y) { return x * y; };
31 std::cout << "Product: " << multiply(3, 4.5) << std::endl;
32
33 // Lambda with STL algorithms
34 std::vector<int> numbers = {1, 2, 3, 4, 5};
35 std::for_each(numbers.begin(), numbers.end(), [](int n) {
36     std::cout << n << " "; });
37 std::cout << std::endl;
38
39 // C++23: Deducing 'this' in lambdas
40 struct Counter {
41     int count = 0;
42     auto increment() {
43         return [this]() {
44             ++count;
45             std::cout << "Count: " << count << std::endl;
46         };
47     };
48
49     Counter counter;
50     auto inc = counter.increment();
51     inc();
52     inc();
53
54     return 0;
55 }

```

**capture** comma-separated list of variables which are captured/modified by the lambda. Captures cannot have same name as input parameters.

Capture list

- \* & = capture all used variables by reference
- \* = = capture all used variables by copy
- \* varName = by-copy
- \* varName... = by-copy pack-expansion
- \* varName initializer = by-copy w/ initializer
- \* &varName = by-reference
- \* &varName... = by-reference pack-expansion
- \* &varName initializer = by-reference w/ initializer
- \* this = by-reference capture of current object
- \* \*this = by-copy capture of current object
- \* ... = by-copy capture of all objects w/ pack expansion
- \* &... initializer = by-reference w/ initializer and pack expansion

```

1 // If the capture-default is &, subsequent simple captures
  must not begin with &.
2 [&] {}; // OK: by-reference capture default
3 [&, i] {}; // OK: by-reference capture, except i is
  captured by copy
4 [&, &i] {}; // Error: by-reference capture when by-
  reference is the default
5 [&, this] {}; // OK, equivalent to [&]
6 [&, this, i] {}; // OK, equivalent to [&, i]

```

```

1 // If the capture-default is =, subsequent simple captures
  must begin with & or be *this(since C++17) or this(since C
  ++20) .
2 [=] {}; // OK: by-copy capture default
3 [=, &i] {}; // OK: by-copy capture, except i is captured by
  reference
4 [=, *this] {}; // until C++17: Error: invalid syntax
5 // since C++17: OK: captures the enclosing S2
  by copy
6 [=, this] {}; // until C++20: Error: this when = is the
  default
7 // since C++20: OK, same as [=]

```

- **mutex**
- **override**
- **random**

```

1 #include <stdlib.h>
2 int rand(); // returns integer in [0, RAND_MAX]

```

```

1 #include <random>
2 // default_random_engine
3 // philox4x64 -> philox_engine
4 // random_device = non-deterministic generator based on
  hardware entropy
5 std::random_device rd;
6 rd.entropy(); // estimate of random number device entropy.
  Deterministic entropy = 0.
7 std::uniform_real_distribution<double> dist(0.0, 1.0);

```

Distribution list

- \* uniform
  - int
  - real (double)
- \* bernoulli
  - bernoulli
  - binomial
  - negative binomial
  - geometric
- \* Poisson
  - poisson
  - exponential
  - gamma
  - weibull
  - extreme\_value
- \* Normal

- normal
- lognormal
- chi\_squared
- cauchy
- fisher\_f
- student\_t
- \* Sampling
  - discrete
  - piecewise\_constant
  - piecewise\_linear
  - item4
- **range-based for**
- **thread**
- **trailing return type** auto main() --> int {return 0;}

- C++14

- **Variable Templates**
- **Generic Lambdas**

- C++17

- **tuple**

```

1 #include <iostream>
2 #include <string>
3 #include <tuple>
4
5 int main() {
6     // Create a tuple with different types
7     std::tuple<int, std::string, double> person = std::
        make_tuple(25, "Alice", 68.5);
8
9     // Access elements of the tuple using std::get
10    int age = std::get<0>(person);
11    std::string name = std::get<1>(person);
12    double weight = std::get<2>(person);
13
14    // Modify elements of the tuple
15    std::get<0>(person) = 30;
16    std::get<2>(person) = 70.0;
17    // Use std::tie to unpack tuple into variables
18    int newAge;
19    std::string newName;
20    double newWeight;
21    std::tie(newAge, newName, newWeight) = person;
22
23    std::cout << "Unpacked Name: " << newName << ", Unpacked
        Age: " << newAge << ", Unpacked Weight: " << newWeight
24        << std::endl;
25
26    // Use std::ignore to unpack only specific elements
27    std::tie(std::ignore, newName, std::ignore) = person;
28    std::cout << "Unpacked Name with ignore: " << newName <<
        std::endl;
29
30    return 0;
31 }

```

- **execution policies**

**seq** used to disambiguate parallel algorithm overloading and require that a parallel algorithm's execution must be sequential. This is used by default when no execution policy is specified.

**par** Indicates that a parallel algorithm MAY be parallelized. Synchronization techniques (e.g. mutexes) may be used.

**par\_unseq** A parallel algorithm MAY be parallelized, vectorized, and moved between threads. Vectorization MUST not use any vectorization-unsafe operations (e.g. mutexes and `std::atomic`)

**unseq** An algorithm's execution MAY be vectorized. Synchronization techniques MUST NOT be used. Since C++20 (the rest of the policies were introduced in C++17).

- C++20

- **Modules**

- **Coroutines**

- **Ranges**

Extension/Generalization of algorithms and iterator libraries to make them less error-prone. Ranges are an abstraction of the following:

```
* [begin, end) iterator pair : ranges::sort()
* begin + [0, size) : views::counted()
* [begin, predicate) : views::take_while() (conditionally-terminated sequences)
* [begin, ..) : unbounded (e.g. views::iota())
```

`std::views` // shorthand for `std::ranges::views` TODO: do more usage/investigation on these

- **Midpoint**

Can be used on any arithmetic type, excluding `bool`. Can be used on objects as long as they are not incomplete types. Returns half the sum of the two inputs, no overflow occurs (this is the main reason to use STL rather than custom implementation). Inputs must point to elements in same object, else behavior is undefined. In case of decimal in average, rounds down.

- **using enum**

- **constexpr**

- **string formatting**

- **template concepts**

- **coroutines**

- **modules**

- C++23

- **print/println**

```
1 #include <print>
2 std::print("{0} {2}{1}!", "Hello", 23, "C++");
3 std::println(); // adds newline to std::print();
```

- **byteswap**

```
1 #include <bit>
2 std::byteswap(T n) noexcept; // T can be any integer value
```

- **flat\_map/flat\_set**

## 4.5 Concepts

- **Types**

- **RAII**

**RAII** Resource Acquisition Is Initialization

- **item3**

- **item4**